

Swin Transformer / Video Swin Transformer : Hierarchical Vision Transformer using Shifted Windows

University Of Seoul.
Machine Learning & Artificial Intelligence
Lab.

Junhyuk So
Undergraduate Student

Swin Transformer: Hierarchical Vision Transformer using Shifted Windows

Ze Liu et al., ICCV 2021, Best Paper Award

🏆 State of the Art	Object Detection on COCO test-dev (using additional training data)	
🏆 State of the Art	Instance Segmentation on COCO test-dev (using additional training data)	
🏆 State of the Art	Object Detection on COCO minival (using additional training data)	
🏆 State of the Art	Instance Segmentation on COCO minival (using additional training data)	
🏆 Ranked #8	Semantic Segmentation on ADE20K (using additional training data)	🏆 Ranked #9 Semantic Segmentation on ADE20K val
🏆 State of the Art	Action Recognition on Something-Something V2 (using additional training data)	
🏆 Ranked #2	Action Classification on Kinetics-400 (using additional training data)	
🏆 Ranked #2	Action Classification on Kinetics-600 (using additional training data)	

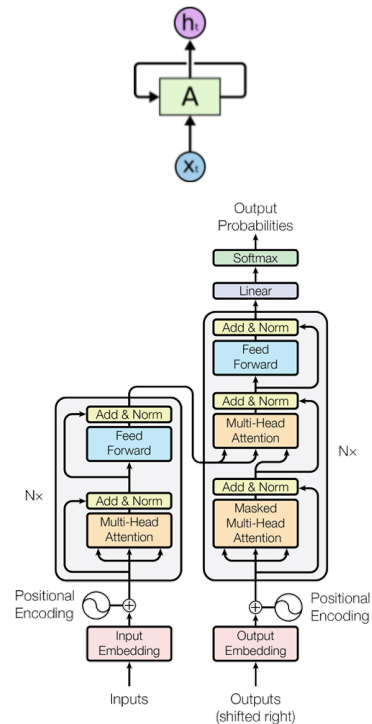
Outline

- Background
 - Transformer
 - Visual Transformer
- Swin Transformer
 - Motivation
 - Method
 - Experiments
- Video Swin Transformer
 - Method
 - Experiments
- Conclusion

Background

Transformer

- Transformer is introduced to solve problems of recurrent model
 - Long term dependency
 - Non-parallel structure
- By Using,
 - (self) Attention Mechanism
 - Multi-head Attention
 - Position Wise Feed Forward
 - Positional Encoding



Background

Attention

- To decide where to **Attend**

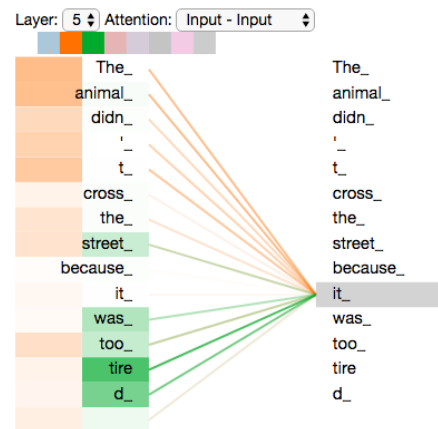
- 1. Transformation of layer map

$$Q = XW^Q, \quad K = YW^K, \quad V = YW^V,$$

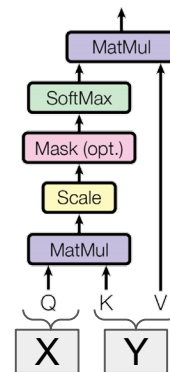
- if self-attention, $X=Y$

- 2. attention layer

- Compare Every Q to Every K
 - by simple dot-product
- Normalize and Softmax
 - Softmax with Q-dimension
- Assign them(Multiply) to Value



Scaled Dot-Product Attention



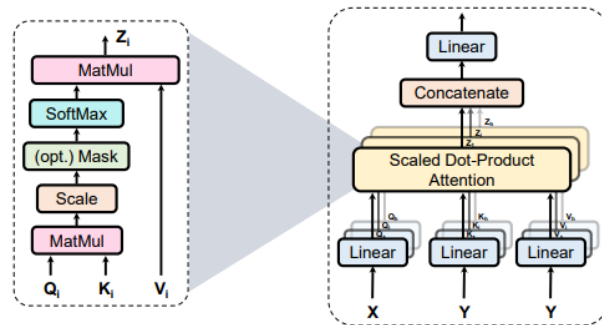
$$\text{Attention}(Q, K, V) = \text{Softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V,$$

Background

Multi Head Attention

- Single Head Attention may have restricted feature subspace
- Solution
 - Use Multiple Feature Subspace by Independent Heads

$$\begin{aligned} Q_i &= XW^{Q_i}, K_i = XW^{K_i}, V_i = XW^{V_i}, \\ Z_i &= \text{Attention}(Q_i, K_i, V_i), i = 1 \dots h, \\ \text{MultiHead}(Q, K, V) &= \text{Concat}(Z_1, Z_2, \dots, Z_h)W^O, \end{aligned} \quad (3)$$



Background

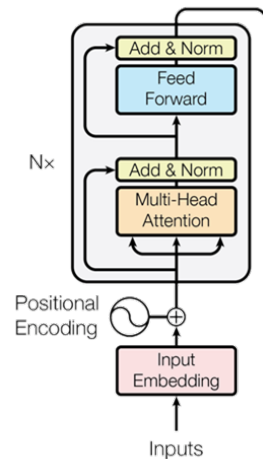
Position Wise Feed Forward Network

- The Output of MHSA is then fed to Two Successive FFN

$$\text{FFN}(x) = \text{RELU}(W_1x + b_1)W_2 + b_2.$$

- Position Wise ?

- [Batch x N_seq x seq_emb1] \rightarrow [Batch x N_seq x seq_emb2]
- 512 \rightarrow 2048 \rightarrow 512



Background

Positional Encoding

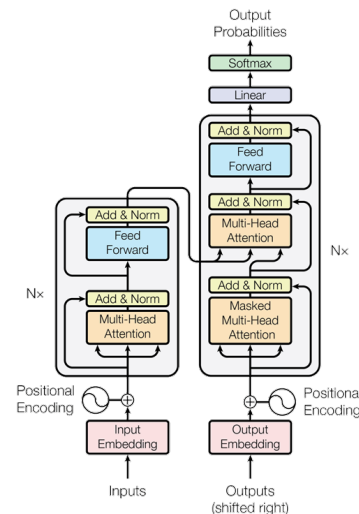
- Attention Mechanism handle All input sequence **identically**
→ the **order** of sequence is **neglected**

- Solution

- add some positional information to input embedding

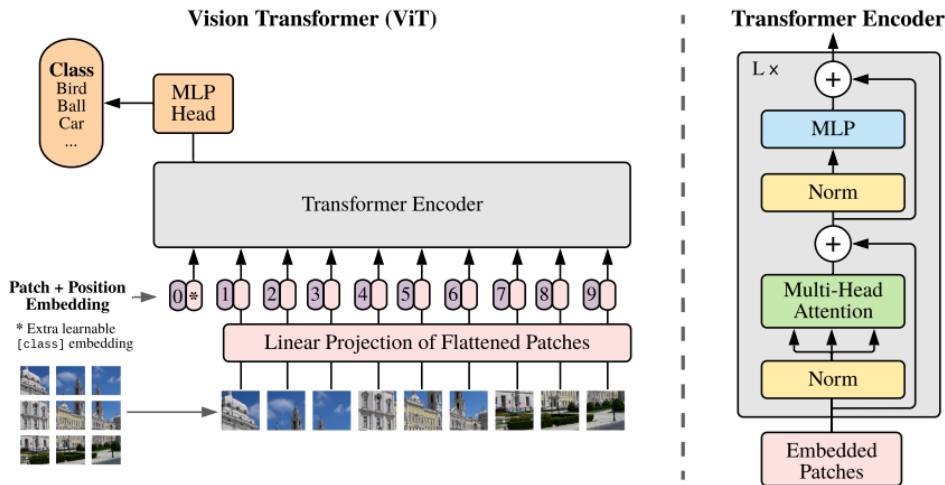
$$PE_{(pos,i)} = \begin{cases} \sin(pos \cdot \omega_k) & \text{if } i = 2k \\ \cos(pos \cdot \omega_k) & \text{if } i = 2k + 1, \end{cases}$$
$$\omega_k = \frac{1}{10000^{2k/d}}, \quad k = 1, \dots, d/2,$$

- where i,d are index, length



Background

Vision Transformer



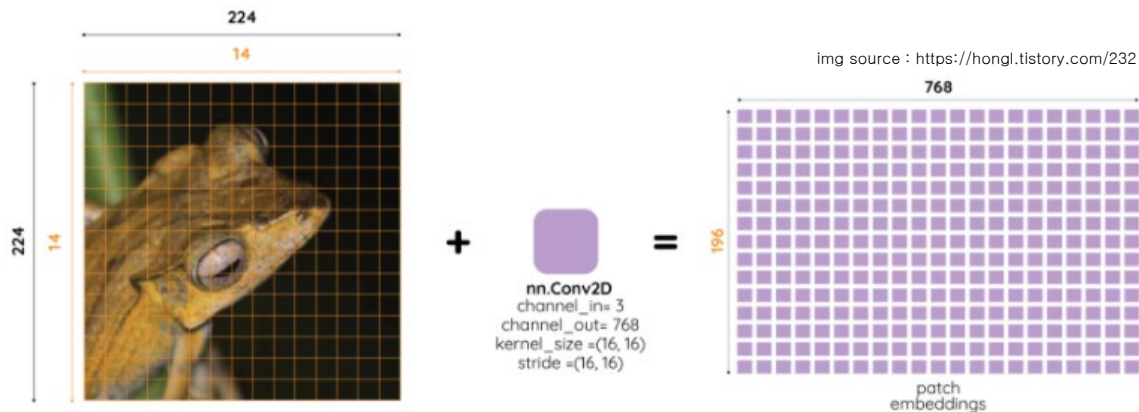
● Almost Same Architecture

- Work Token \rightarrow Image Patch & Linear Projection
- Normalization first

Background

ViT (Vision Transformer)

- Patch Partitioning & Linear Projection
 - $H \times W \times C \rightarrow (H/16) \times (W/16) \times 3$
 - Mix accross channel dim.
 - Can be implmeneted by simple Conv2D.



Swin Transformer

Swin Transformer

Motivation

- problems of original ViT
 - No consideration for Vision domain.
 - Scale-Problem
 - Computational Inefficiency
 - Quadratic Computation
 - Memory Inefficiency

Scale problem

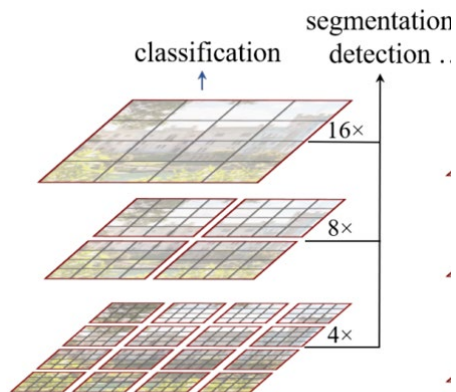
- Size of object can vary depending on Location of Camera
 - Same Patch size ignores this issue.



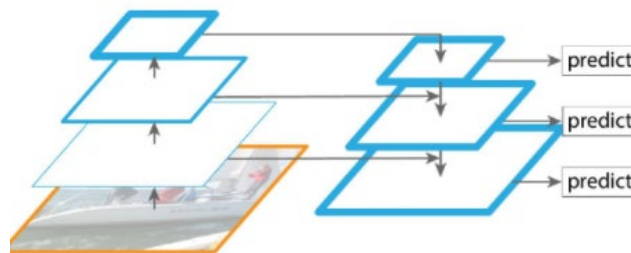
Scale problem

Solution

- Use Different Patch Size!
- Aggregate differently scaled feature map to make final prediction.
 - E.g) FPN



(a) Swin Transformer (ours)



(d) Feature Pyramid Network

Quadratic Computation

- Attention Mechanism → Quadratic Complexity

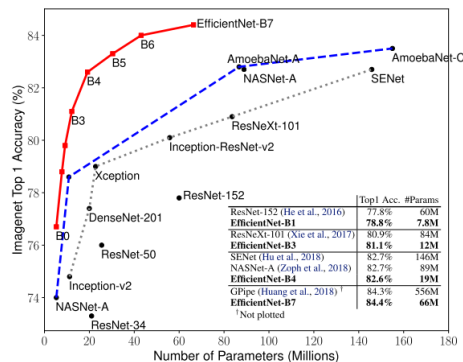
- hard to use larger image size (resolution)

$$A(Q, K, V) = \text{softmax}(QK^T)V$$

- But why we need larger image Size ?

- Vision tasks tend to generalize well when use bigger image size

■ E.g) EfficientNet B0 : 224, EfficientNetB7 : 600



Quadratic Computation

● Solution

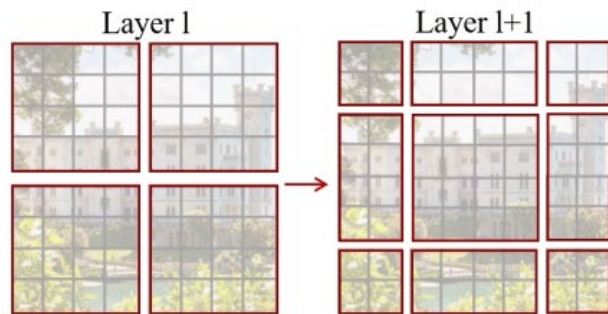
- Local Attention (W – MSA)
 - Attend only in windows

● Can have **Linear** computational complexity :)

● Can add inductive bias :)?

● Can't handle **long range**-dependancy :(

- Use Shifted Window (SW-MSA)



$$\Omega(\text{MSA}) = 4hwC^2 + 2(hw)^2C,$$

$$\Omega(\text{W-MSA}) = 4hwC^2 + 2M^2hwC,$$

Overall Architecture

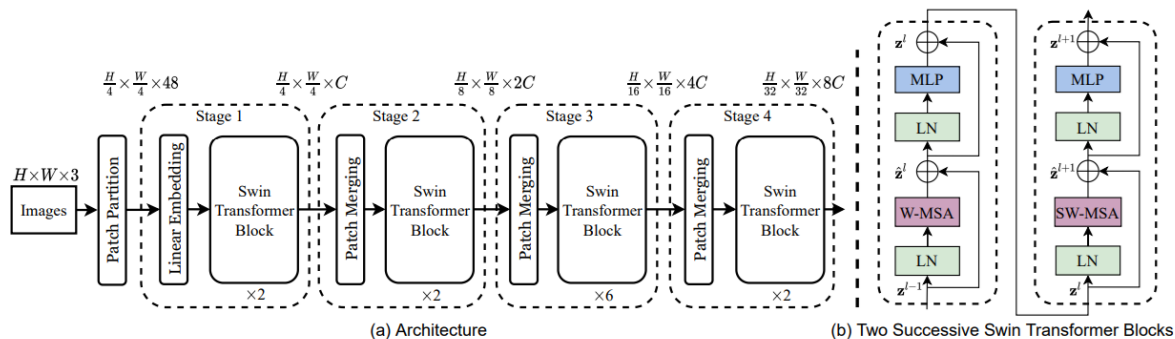
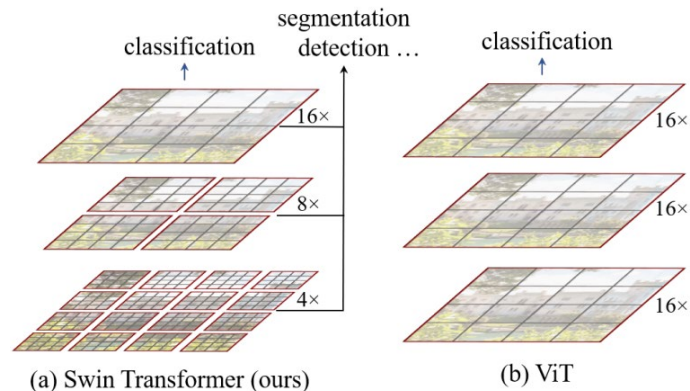


Figure 3. (a) The architecture of a Swin Transformer (Swin-T); (b) two successive Swin Transformer Blocks (notation presented with Eq. (3)). W-MSA and SW-MSA are multi-head self attention modules with regular and shifted windowing configurations, respectively.



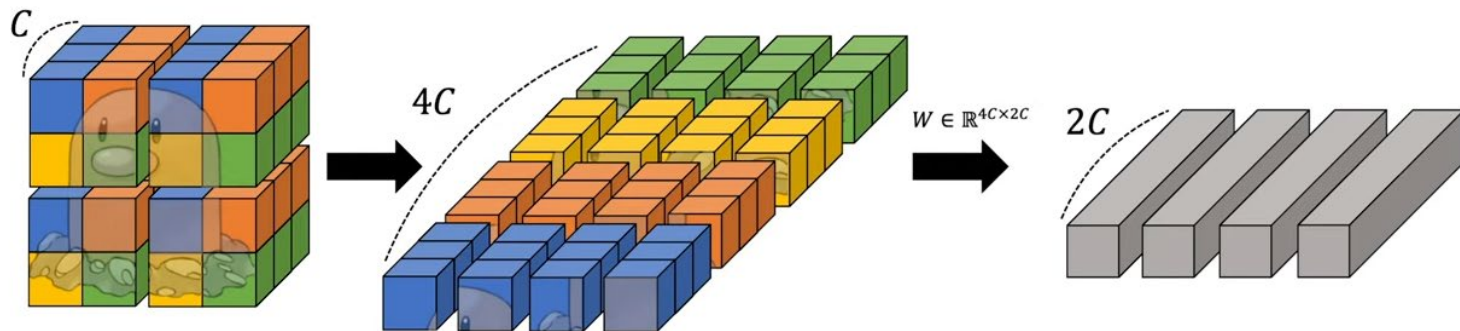
Details

- Patch Merging Block
- Relative Positional Bias
- Cyclic Shifted Batch Computing

Method

Patch Merging Block

- If we simply merge 2x2 Neighbor patches..
 - $1C \rightarrow 4C$
- Use Linear projection across Channel dimension..
 - $4C \rightarrow 2C$



Method

Relative Position Bias

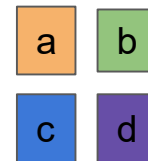
Relative position bias In computing self-attention, we follow [49, 1, 32, 33] by including a relative position bias $B \in \mathbb{R}^{M^2 \times M^2}$ to each head in computing similarity:

$$\text{Attention}(Q, K, V) = \text{SoftMax}(QK^T / \sqrt{d} + B)V, \quad (4)$$

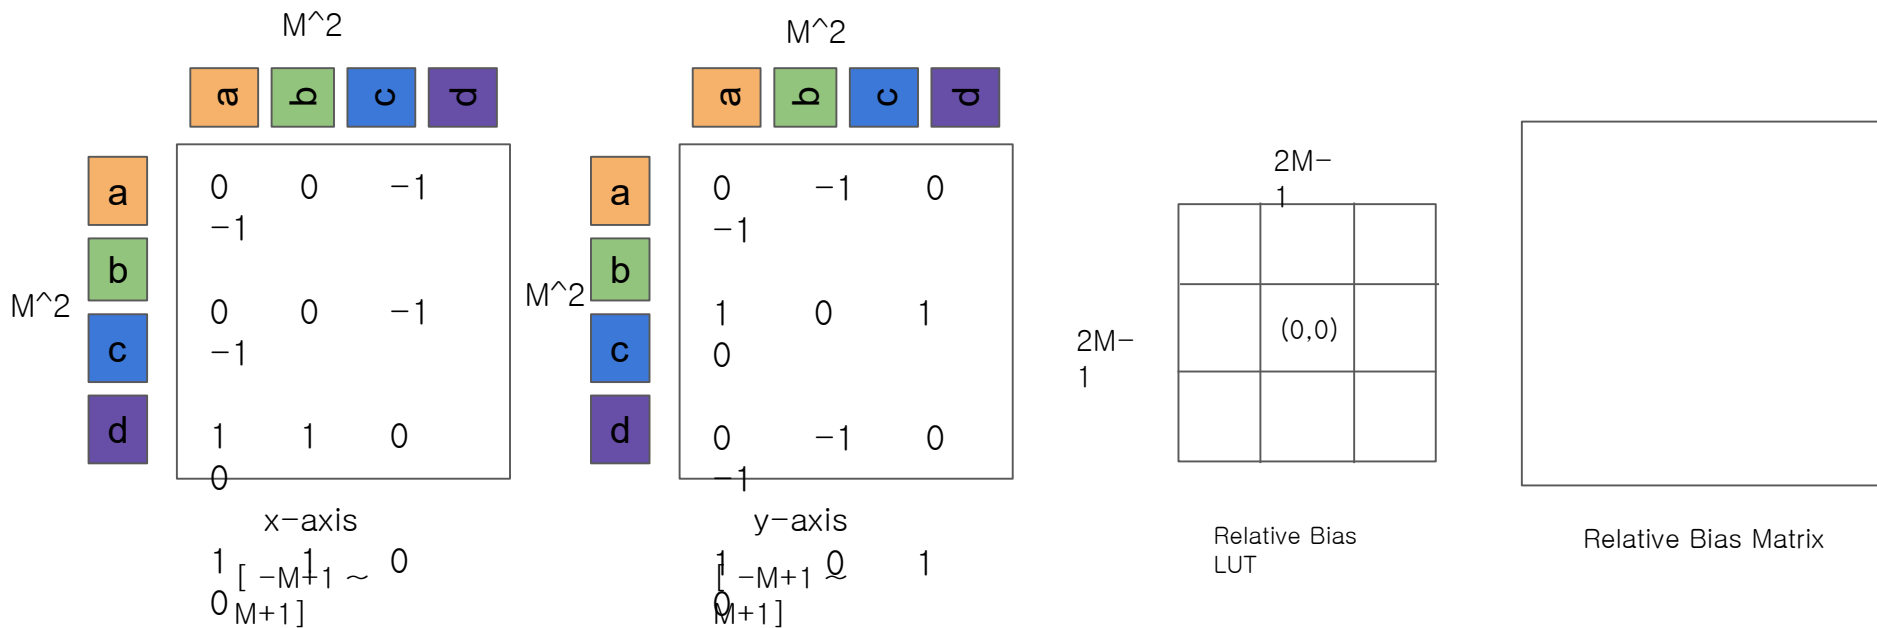
where $Q, K, V \in \mathbb{R}^{M^2 \times d}$ are the *query*, *key* and *value* matrices; d is the *query/key* dimension, and M^2 is the number of patches in a window. Since the relative position along each axis lies in the range $[-M + 1, M - 1]$, we parameterize a smaller-sized bias matrix $\hat{B} \in \mathbb{R}^{(2M-1) \times (2M-1)}$, and values in B are taken from \hat{B} .

Method

$M=2$

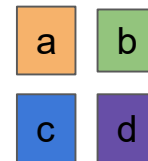


Relative Position Bias

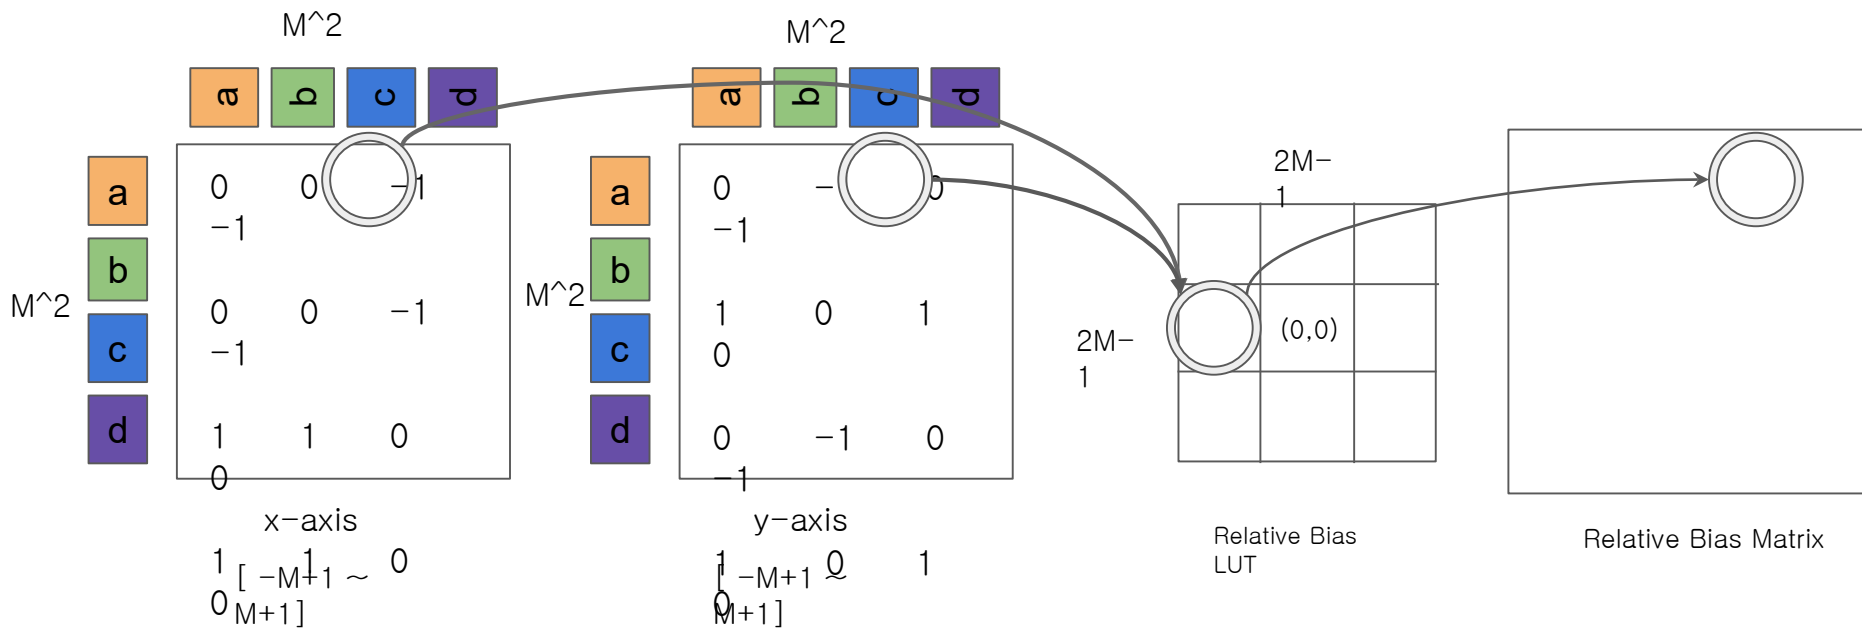


Method

$M=2$



Relative Position Bias



Cyclic Shifted Batch Computing

- Problem : If we use SW-MSA, Window sizes are not same
 - Naive solution : Pad every small windows
 - But this make # Patch $N+1, N+1$. \rightarrow More Computation !
- Solution : Cyclic Shift and Masked MSA

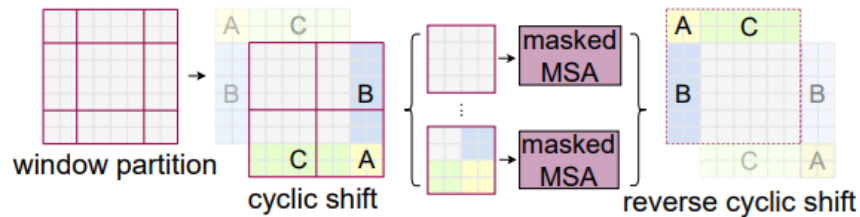


Figure 4. Illustration of an efficient batch computation approach for self-attention in shifted window partitioning.

Cyclic Shifted Batch Computing

Masked
MSA

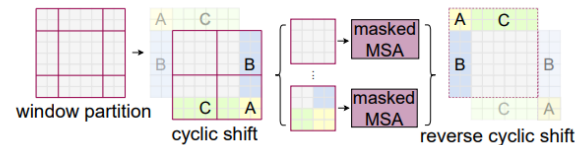
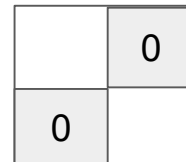
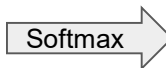
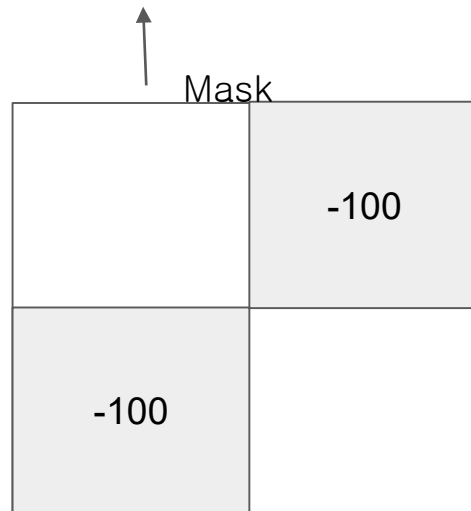
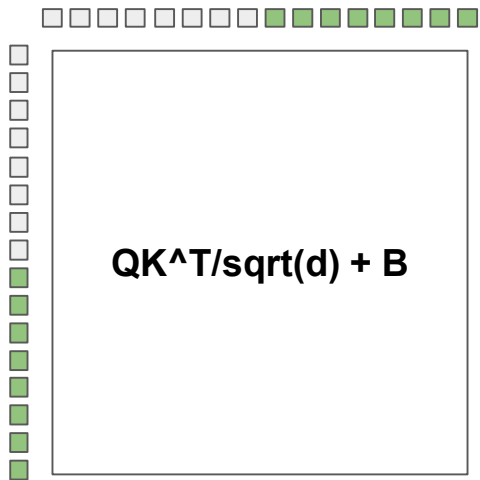
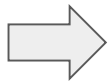
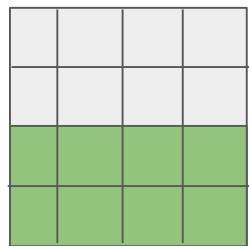


Figure 4. Illustration of an efficient batch computation approach for self-attention in shifted window partitioning.

Experiments

Datasets

- Classification
 - Imagenet 1k
 - Imagenet 1k – ImageNet 22K Pretrained
- Detection
 - COCO 2017
- Segmentation
 - ADE20K

Settings

- Optimizer , LR
 - AdamW ($1e-3$, bs 1024)
- Augmentation
 - RandAug, Mixup, Cutmix, Cutout
 - Stochastic Depth.
 - No Repeated Aug.
- Regularization
 - Weight decay, Gradient clipping

Experiments

Shifted Windows w/o, Positional Embedding w/po

	ImageNet		COCO		ADE20k
	top-1	top-5	AP ^{box}	AP ^{mask}	mIoU
w/o shifting	80.2	95.1	47.7	41.5	43.3
shifted windows	81.3	95.6	50.5	43.7	46.1
no pos.	80.1	94.9	49.2	42.6	43.8
abs. pos.	80.5	95.2	49.0	42.4	43.2
abs.+rel. pos.	81.3	95.6	50.2	43.4	44.0
rel. pos. w/o app.	79.3	94.7	48.2	41.9	44.1
rel. pos.	81.3	95.6	50.5	43.7	46.1

Table 4. Ablation study on the *shifted windows* approach and different position embedding methods on three benchmarks, using the Swin-T architecture. w/o shifting: all self-attention modules adopt regular window partitioning, without *shifting*; abs. pos.: absolute position embedding term of ViT; rel. pos.: the default settings with an additional relative position bias term (see Eq. (4)); app.: the first scaled dot-product term in Eq. (4).

Experiments

Imagenet Classification on ImageNet-1k

(a) Regular ImageNet-1K trained models					
method	image size	#param.	FLOPs	throughput (image / s)	ImageNet top-1 acc.
RegNetY-4G [48]	224 ²	21M	4.0G	1156.7	80.0
RegNetY-8G [48]	224 ²	39M	8.0G	591.6	81.7
RegNetY-16G [48]	224 ²	84M	16.0G	334.7	82.9
EffNet-B3 [58]	300 ²	12M	1.8G	732.1	81.6
EffNet-B4 [58]	380 ²	19M	4.2G	349.4	82.9
EffNet-B5 [58]	456 ²	30M	9.9G	169.1	83.6
EffNet-B6 [58]	528 ²	43M	19.0G	96.9	84.0
EffNet-B7 [58]	600 ²	66M	37.0G	55.1	84.3
ViT-B/16 [20]	384 ²	86M	55.4G	85.9	77.9
ViT-L/16 [20]	384 ²	307M	190.7G	27.3	76.5
DeiT-S [63]	224 ²	22M	4.6G	940.4	79.8
DeiT-B [63]	224 ²	86M	17.5G	292.3	81.8
DeiT-B [63]	384 ²	86M	55.4G	85.9	83.1
Swin-T	224 ²	29M	4.5G	755.2	81.3
Swin-S	224 ²	50M	8.7G	436.9	83.0
Swin-B	224 ²	88M	15.4G	278.1	83.5
Swin-B	384 ²	88M	47.0G	84.7	84.5

(b) ImageNet-22K pre-trained models					
method	image size	#param.	FLOPs	throughput (image / s)	ImageNet top-1 acc.
R-101x3 [38]	384 ²	388M	204.6G	-	84.4
R-152x4 [38]	480 ²	937M	840.5G	-	85.4
ViT-B/16 [20]	384 ²	86M	55.4G	85.9	84.0
ViT-L/16 [20]	384 ²	307M	190.7G	27.3	85.2
Swin-B	224 ²	88M	15.4G	278.1	85.2
Swin-B	384 ²	88M	47.0G	84.7	86.4
Swin-L	384 ²	197M	103.9G	42.1	87.3

Experiments

Object detection on COCO

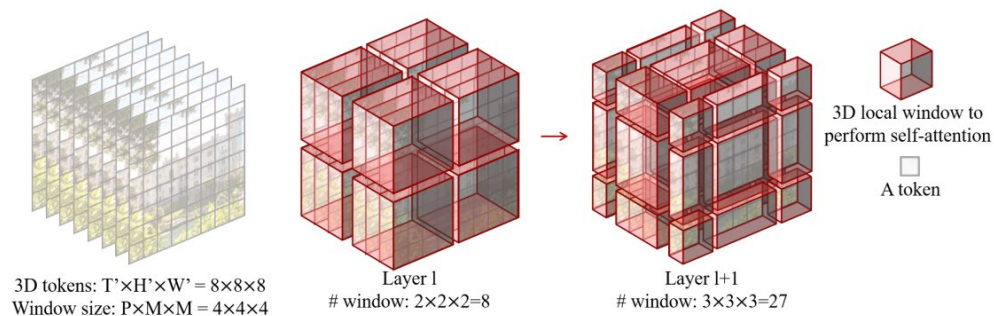
(a) Various frameworks							
Method	Backbone	AP ^{box}	AP ^{box} ₅₀	AP ^{box} ₇₅	#param.	FLOPs	FPS
Cascade	R-50	46.3	64.3	50.5	82M	739G	18.0
Mask R-CNN	Swin-T	50.5	69.3	54.9	86M	745G	15.3
ATSS	R-50	43.5	61.9	47.0	32M	205G	28.3
	Swin-T	47.2	66.5	51.3	36M	215G	22.3
RepPointsV2	R-50	46.5	64.6	50.3	42M	274G	13.6
	Swin-T	50.0	68.5	54.2	45M	283G	12.0
Sparse R-CNN	R-50	44.5	63.4	48.2	106M	166G	21.0
	Swin-T	47.9	67.3	52.3	110M	172G	18.4

(b) Various backbones w. Cascade Mask R-CNN									
	AP ^{box}	AP ^{box} ₅₀	AP ^{box} ₇₅	AP ^{mask}	AP ^{mask} ₅₀	AP ^{mask} ₇₅	#param	FLOPs	FPS
DeiT-S [†]	48.0	67.2	51.7	41.4	64.2	44.3	80M	889G	10.4
R50	46.3	64.3	50.5	40.1	61.7	43.4	82M	739G	18.0
Swin-T	50.5	69.3	54.9	43.7	66.6	47.1	86M	745G	15.3
X101-32	48.1	66.5	52.4	41.6	63.9	45.2	101M	819G	12.8
Swin-S	51.8	70.4	56.3	44.7	67.9	48.5	107M	838G	12.0
X101-64	48.3	66.4	52.3	41.7	64.0	45.1	140M	972G	10.4
Swin-B	51.9	70.9	56.5	45.0	68.4	48.7	145M	982G	11.6

Video Swin Transformer

Ze Liu et al, 2021 Arxiv

Motivation



- Spatio-temporal adaption of Swin Transformer
- strictly follows the hierarchical structure of the original Swin
 - but extends the scope of local attention computation from only the spatial domain to the spatiotemporal domain

Method

Overall Architecture

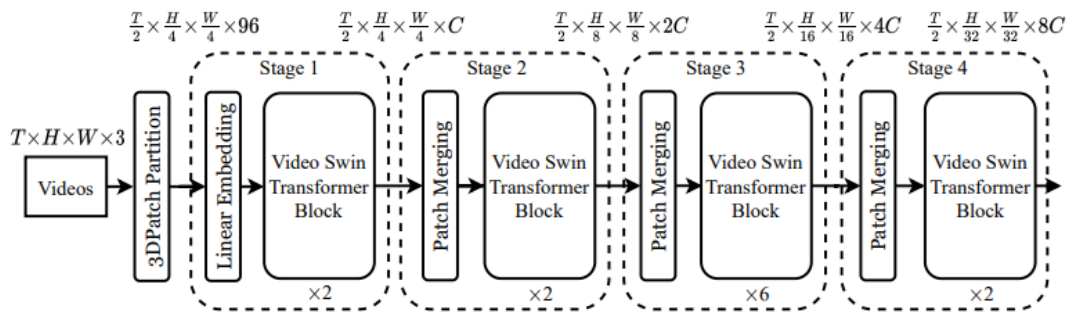
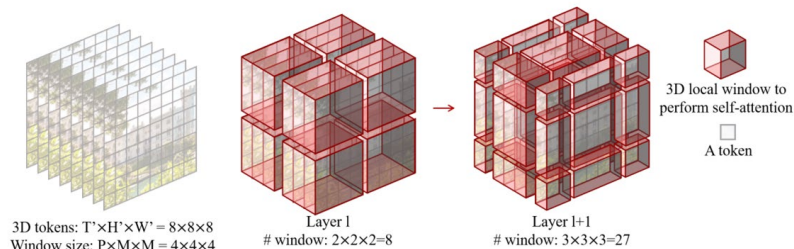


Figure 1: Overall architecture of Video Swin Transformer (tiny version, referred to as Swin-T).



Details

- No patch merging between T-dimension
- 3D Relative Position Bias

Method

3D Relational Position Bias

3D Relative Position Bias Numerous previous works [31, 2, 16, 17] have shown that it can be advantageous to include a relative position bias to each head in self-attention computation. Thus, we follow [28] by introducing 3D relative position bias $B \in \mathbb{R}^{P^2 \times M^2 \times M^2}$ for each head as

$$\text{Attention}(Q, K, V) = \text{SoftMax}(QK^T / \sqrt{d} + B)V, \quad (2)$$

where $Q, K, V \in \mathbb{R}^{PM^2 \times d}$ are the *query*, *key* and *value* matrices; d is the dimension of *query* and *key* features, and PM^2 is the number of tokens in a 3D window. Since the relative position along each axis lies in the range of $[-P + 1, P - 1]$ (temporal) or $[-M + 1, M - 1]$ (height or width), we parameterize a smaller-sized bias matrix $\hat{B} \in \mathbb{R}^{(2P-1) \times (2M-1) \times (2M-1)}$, and values in B are taken from \hat{B} .



3D Bias LUT

Experiments

Table 1: Comparison to state-of-the-art on Kinetics-400. "384↑" signifies that the model uses a large spatial resolution of 384×384 . "Views" indicates # temporal clip \times # spatial crop. The magnitude are Giga (10^9) and Mega (10^6) for FLOPs and Param respectively.

Method	Pretrain	Top-1	Top-5	Views	FLOPs	Param
R(2+1)D [37]	-	72.0	90.0	10×1	75	61.8
I3D [6]	ImageNet-1K	72.1	90.3	-	108	25.0
NL I3D-101 [40]	ImageNet-1K	77.7	93.3	10×3	359	61.8
ip-CSN-152 [36]	-	77.8	92.8	10×3	109	32.8
CorrNet-101 [39]	-	79.2	-	10×3	224	-
SlowFast R101+NL [13]	-	79.8	93.9	10×3	234	59.9
X3D-XXL [12]	-	80.4	94.6	10×3	144	20.3
MViT-B, 32×3 [10]	-	80.2	94.4	1×5	170	36.6
MViT-B, 64×3 [10]	-	81.2	95.1	3×3	455	36.6
TimeSformer-L [3]	ImageNet-21K	80.7	94.7	1×3	2380	121.4
ViT-B-VTN [29]	ImageNet-21K	78.6	93.7	1×1	4218	11.04
ViViT-L/16x2 [1]	ImageNet-21K	80.6	94.7	4×3	1446	310.8
ViViT-L/16x2 320 [1]	ImageNet-21K	81.3	94.7	4×3	3992	310.8
ip-CSN-152 [36]	IG-65M	82.5	95.3	10×3	109	32.8
ViViT-L/16x2 [1]	JFT-300M	82.8	95.5	4×3	1446	310.8
ViViT-L/16x2 320 [1]	JFT-300M	83.5	95.5	4×3	3992	310.8
ViViT-H/16x2 [1]	JFT-300M	84.8	95.8	4×3	8316	647.5
Swin-T	ImageNet-1K	78.8	93.6	4×3	88	28.2
Swin-S	ImageNet-1K	80.6	94.5	4×3	166	49.8
Swin-B	ImageNet-1K	80.6	94.6	4×3	282	88.1
Swin-B	ImageNet-21K	82.7	95.5	4×3	282	88.1
Swin-L	ImageNet-21K	83.1	95.9	4×3	604	197.0
Swin-L (384↑)	ImageNet-21K	84.6	96.5	4×3	2107	200.0
Swin-L (384↑)	ImageNet-21K	84.9	96.7	10×5	2107	200.0

Video Classification

- Pretraining on ImageNet
 - Directly duplicate weights in Swin Twice
 - Multiply whole matrix 0.5

Experiments

Ablation Study

Temporal dimension/window size

Table 5: Ablation study on temporal dimension of 3D tokens and temporal window size with Swin-T on K400.

temporal dimension	Window size	Top 1	Top 5	FLOPs	Param
16	$16 \times 7 \times 7$	79.1	93.8	106	28.5
8	$8 \times 7 \times 7$	78.5	93.2	44	28.2
4	$4 \times 7 \times 7$	76.7	92.5	20	28.0
16	$16 \times 7 \times 7$	79.1	93.8	106	28.5
16	$8 \times 7 \times 7$	78.8	93.6	88	28.2
16	$4 \times 7 \times 7$	78.6	93.4	79	28.0

Temporal shifting Effect

Table 6: Ablation study on the 3D shifted window approach with Swin-T on K400.

	Top-1	Top-5
w. 3D shifting	78.8	93.6
w/o temporal shifting	78.5	93.5
w/o 3D shifting	78.1	93.3

Conclusion

Swin Transformer

- Solve scale problem by Dynamic Patch Size
- Solve Quadratic Time Complexity problem by Local Attention
 - Solve Long range dependency problem by Shifted Window

Future research ?

- Temporal Hierarchy
- Temporal scale problem

Q&A