

2022/03/14

Lab Seminar

Domain Generalization
Gradient Matching for Domain Generalization

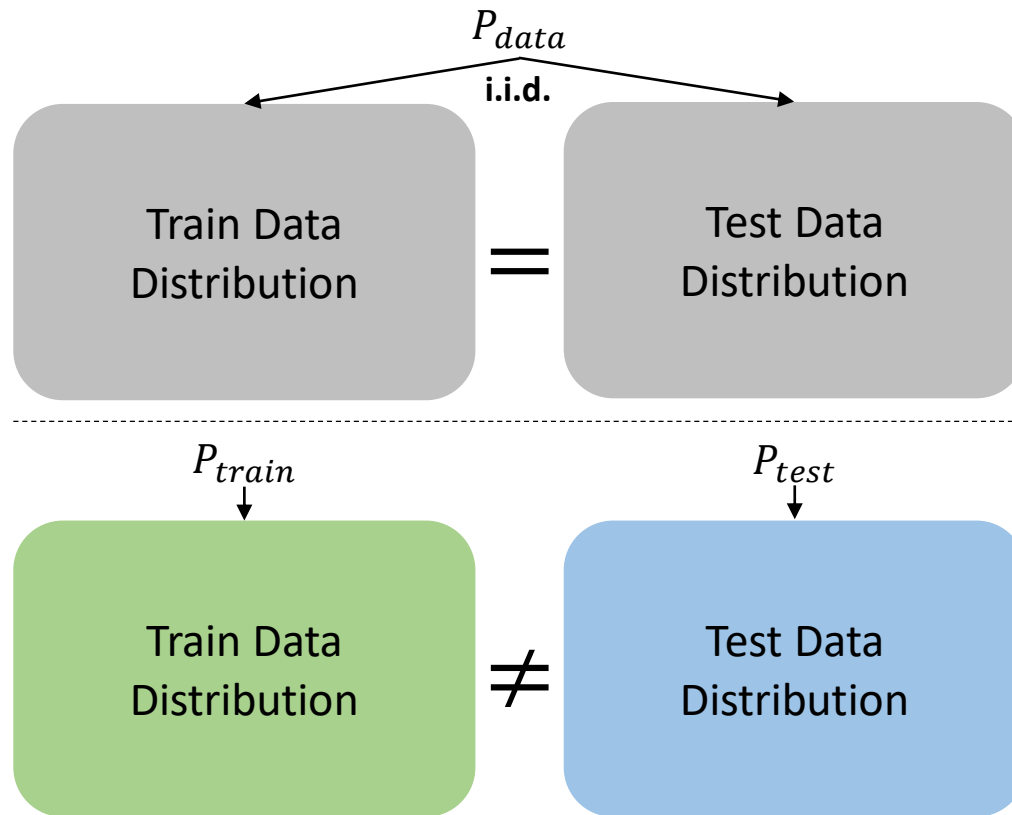
Taero Kim

MLAI, Univ. of Seoul

Domain Generalization

Domain Generalization

- Distribution Shift



We take the assumption for granted that **train data and test data are sampled from same distribution**.

In reality, however, **distributions of train and test data are always different each other** unless the users adjust them carefully. We call this situation **“Distribution Shift”**.

Domain Generalization

• Distribution Shift

There are many types of Distribution Shift;

- Covariate Shift, **Domain Shift**, **Subpopulation Shift**, Continual Shift, etc.


In this seminar, we will consider only about domain shift and subpopulation shift.

To understand **the difference between domain shift and subpopulation shift**, we can introduce Darts game.

Let d be the distance between the player and the dart board.

Less than 10 points = label 0 / More than 10 points = label 1

The player practiced 8 times when $d=3m$ and 2 times when $d=7m$.

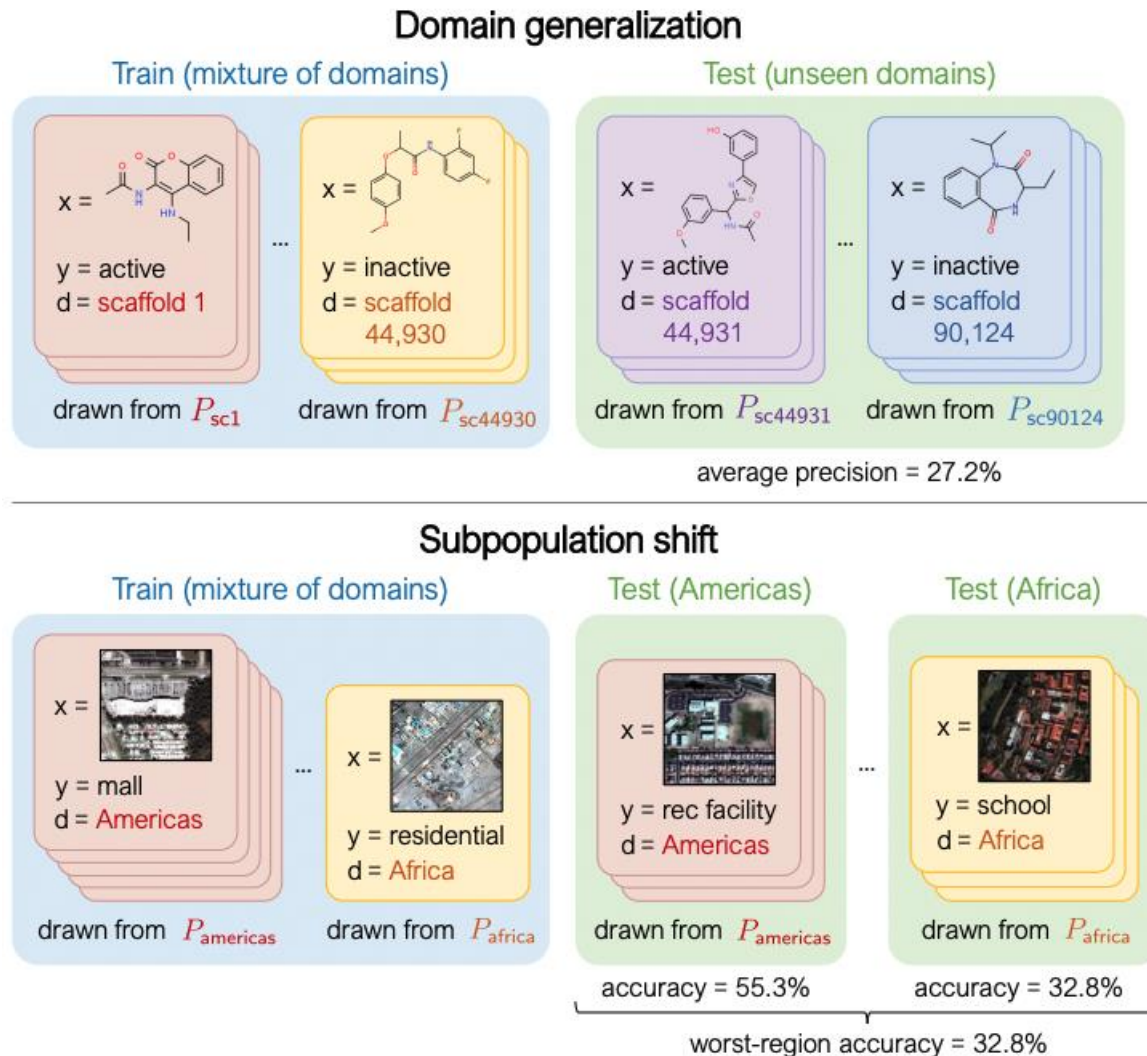


Pictures from Pixabay

Domain Generalization

- Domain Generalization and Subpopulation Shift

Wilds Dataset 1.2v


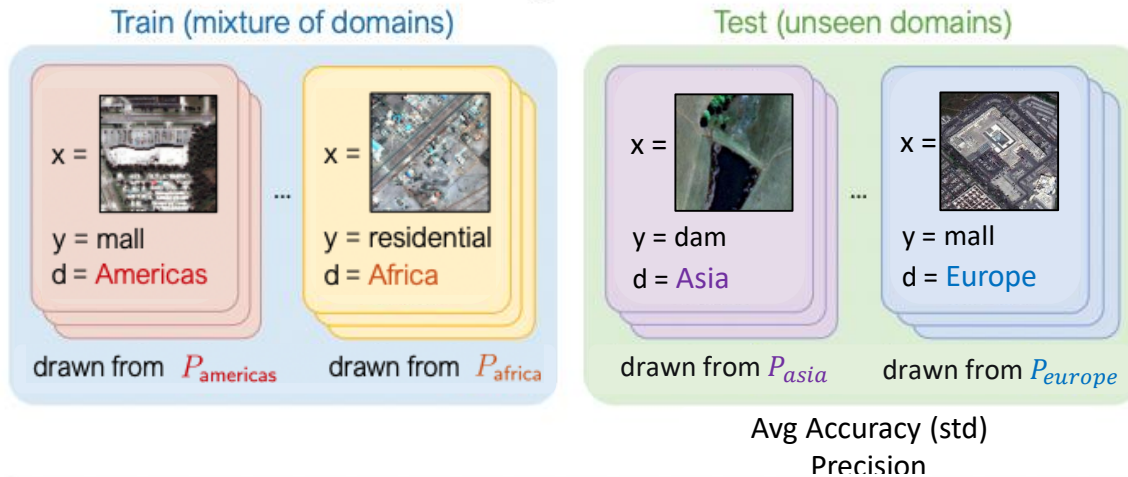


Domain Generalization

Domain Generalization VS Subpopulation Shift

fMoW Datasets in Wilds 1.2v

Domain generalization



Low accuracy
due to low precision

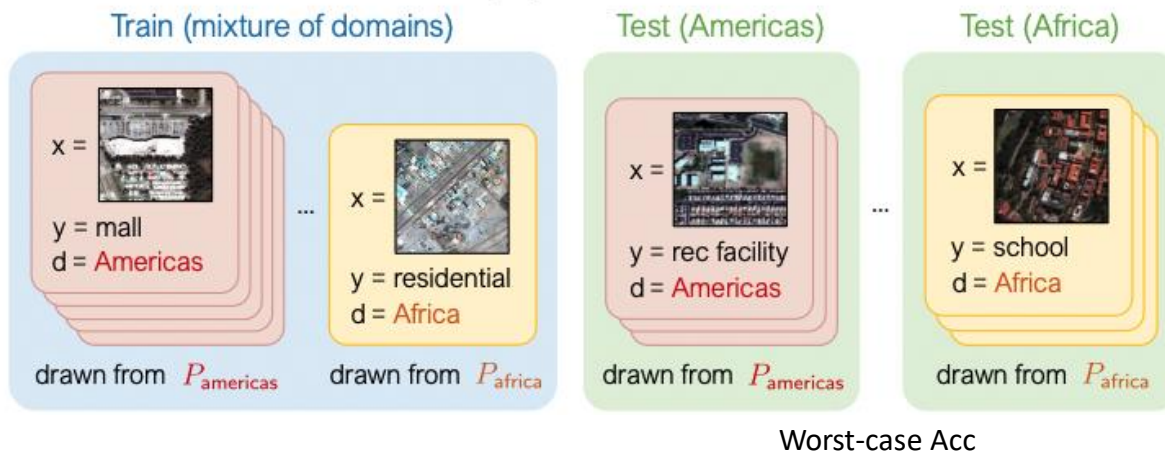


Low accuracy
even with high precision

Goal of DG is to get high performance about related but unseen domain.

The domain can vary depending on how we set it.

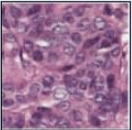
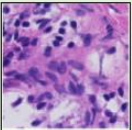
Subpopulation shift




Ensuring Subpopulation Shift is one way to improve DG. Their methodologies are also similar. However, subpopulation shift alone cannot guarantee performance for unseen domain.

Inspired by Christie et al., 2018; Koh et al., 2021; Wikipedia Accuracy and Precision;

- Camelyon17 Dataset – Wilds 1.2

Dataset	Camelyon17
Input (x)	o tissue slide
Prediction (y)	tumor
Domain (d)	hospital
# domains	5
# examples	455,954
Train example	
Test example	
Adapted from	Bandi et al. 2018



450 Lymph Node in human

Domain Generalization does not always guarantee worst-case accuracy.

Center A	Slide 1	...	Slide10
Normal (0)	97%		48%
Tumor (1)	3%		52%

Worst-case

====> Epoch: 019

===== val =====

Average acc: 0.838

slide = 10 [n = 2078]: acc = 0.861

slide = 11 [n = 5371]: acc = 0.490

slide = 12 [n = 1705]: acc = 0.880

slide = 13 [n = 10404]: acc = 0.948

slide = 14 [n = 888]: acc = 0.957

slide = 15 [n = 2966]: acc = 0.973

slide = 16 [n = 4971]: acc = 0.869

slide = 17 [n = 2659]: acc = 0.744

slide = 18 [n = 1275]: acc = 0.932

slide = 19 [n = 2587]: acc = 0.866

Center A

Slides

Worst-group acc: 0.490

===== test =====

Average acc: 0.802

slide = 20 [n = 3810]: acc = 0.580

slide = 21 [n = 3694]: acc = 0.469

slide = 22 [n = 7210]: acc = 0.856

slide = 23 [n = 5288]: acc = 0.840

slide = 24 [n = 7727]: acc = 0.728

slide = 25 [n = 4334]: acc = 0.770

slide = 26 [n = 3815]: acc = 0.693

slide = 27 [n = 4556]: acc = 0.536


slide = 28 [n = 31878]: acc = 0.917

slide = 29 [n = 12742]: acc = 0.813

Center B

Slides

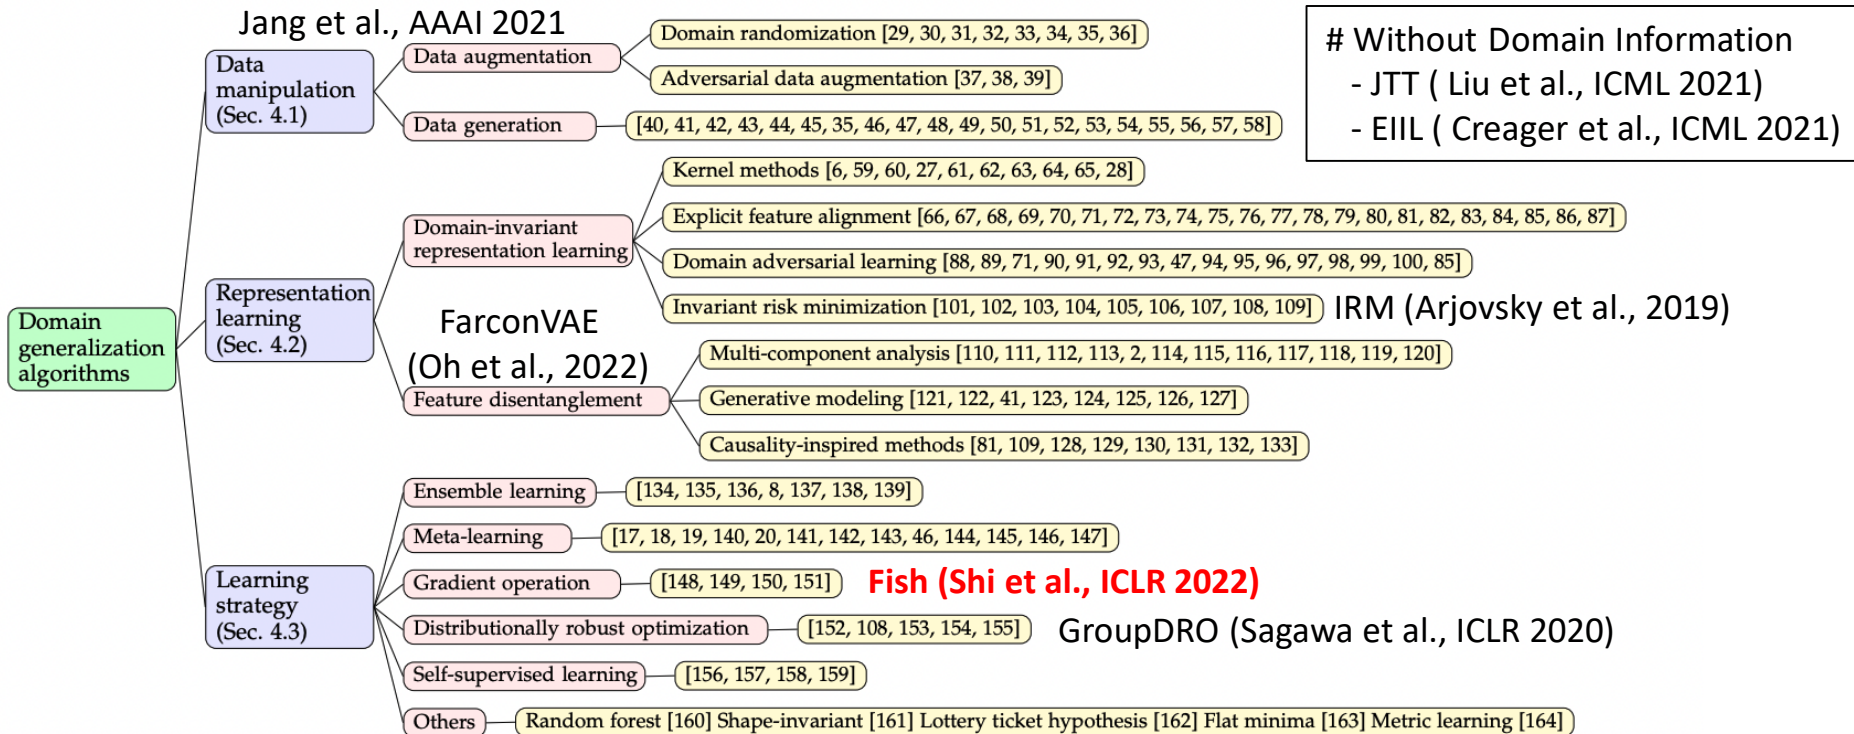
Worst-group acc: 0.469



Domain Generalization

Algorithms Taxonomy for Domain Generalization

Sorted by method

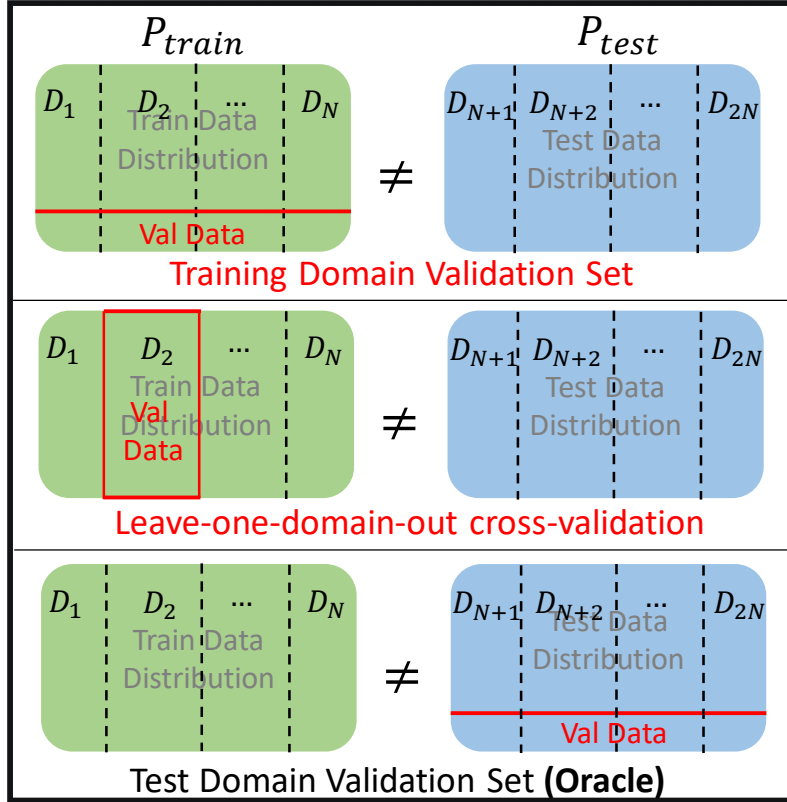


This picture shows the taxonomy of domain generalization algorithms. DG is a topic that has recently been in the spotlight, and it is not determined which methods works better.

Domain Generalization

• DomainBed – Model Selection Method & DG Benchmark

Recommendation 1 A domain generalization algorithm should be responsible for specifying a model selection method.



Recommendation 2 Researchers should disclaim any oracle-selection results as such and specify policies to limit access to the test domain.

$$\text{ERM Loss } L_{ERM}(\mathcal{D}_{tr}; \theta) = \mathbb{E}_{\mathcal{D} \sim \mathcal{D}_{tr}} \mathbb{E}_{(x,y) \sim \mathcal{D}} [\ell((x,y); \theta)]$$

ERM is powerful algorithms for Domain Generalization

Model selection method: training domain validation set

Algorithm	CMNIST	RMNIST	VLCS	PACS	Office-Home	TerraInc	DomainNet	Avg
ERM	52.0 ± 0.1	98.0 ± 0.0	77.4 ± 0.3	85.7 ± 0.5	67.5 ± 0.5	47.2 ± 0.4	41.2 ± 0.2	67.0
IRM	51.8 ± 0.1	97.9 ± 0.0	78.1 ± 0.0	84.4 ± 1.1	66.6 ± 1.0	47.9 ± 0.7	35.7 ± 1.9	66.0
DRO	52.0 ± 0.1	98.1 ± 0.0	77.2 ± 0.6	84.1 ± 0.4	66.9 ± 0.3	47.0 ± 0.3	33.7 ± 0.2	65.5
Mixup	51.9 ± 0.1	98.1 ± 0.0	77.7 ± 0.4	84.3 ± 0.5	69.0 ± 0.1	48.9 ± 0.8	39.6 ± 0.1	67.1
MLDG	51.6 ± 0.1	98.0 ± 0.0	77.1 ± 0.4	84.8 ± 0.6	68.2 ± 0.1	46.1 ± 0.8	41.8 ± 0.4	66.8
CORAL	51.7 ± 0.1	98.1 ± 0.1	77.7 ± 0.5	86.0 ± 0.2	68.6 ± 0.4	46.4 ± 0.8	41.8 ± 0.2	67.2
MMD	51.8 ± 0.1	98.1 ± 0.0	76.7 ± 0.9	85.0 ± 0.2	67.7 ± 0.1	49.3 ± 1.4	39.4 ± 0.8	66.8
DANN	51.5 ± 0.3	97.9 ± 0.1	78.7 ± 0.3	84.6 ± 1.1	65.4 ± 0.6	48.4 ± 0.5	38.4 ± 0.0	66.4
C-DANN	51.9 ± 0.1	98.0 ± 0.0	78.2 ± 0.4	82.8 ± 1.5	65.6 ± 0.5	47.6 ± 0.8	38.9 ± 0.1	66.1

Model selection method: Leave-one-domain-out cross-validation

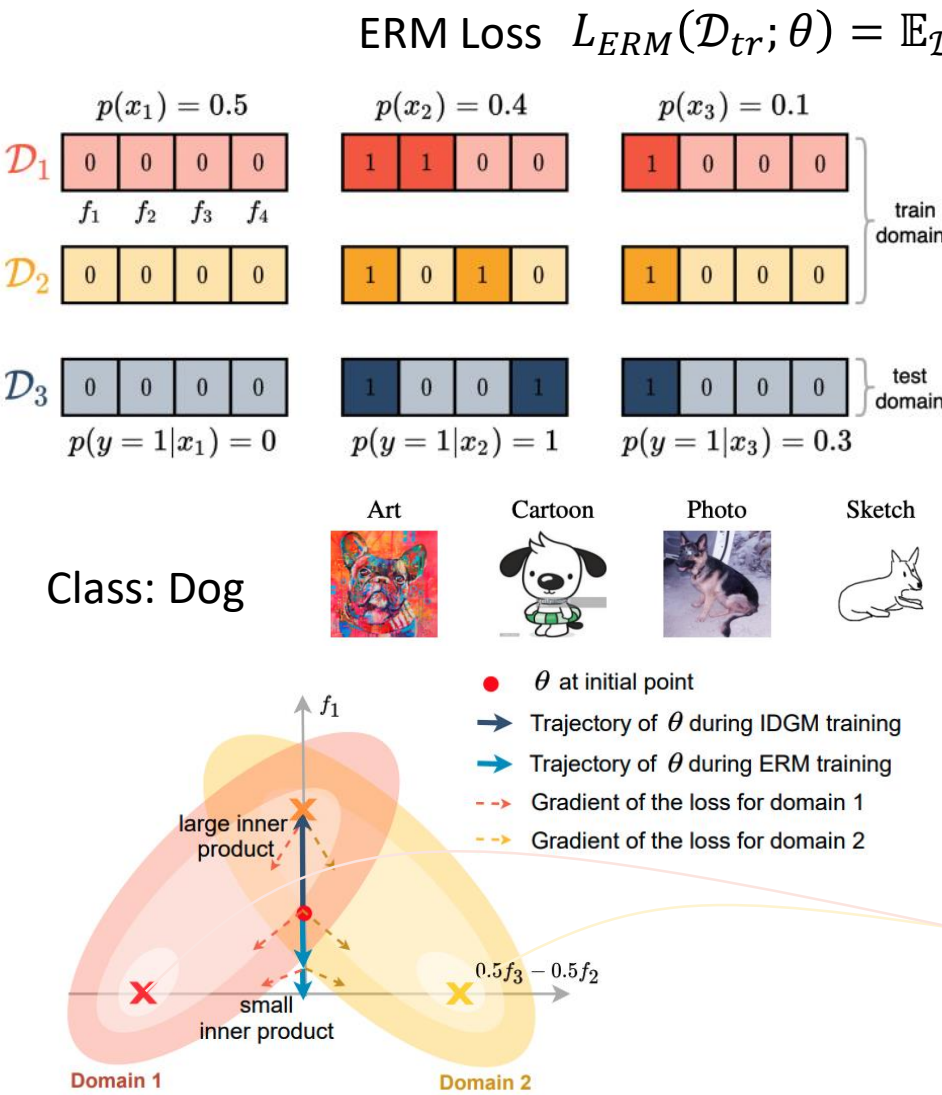
Algorithm	CMNIST	RMNIST	VLCS	PACS	Office-Home	TerraInc	DomainNet	Avg
ERM	34.2 ± 1.2	98.0 ± 0.0	76.8 ± 1.0	83.3 ± 0.6	67.3 ± 0.3	46.2 ± 0.2	40.8 ± 0.2	63.8
IRM	36.3 ± 0.4	97.7 ± 0.1	77.2 ± 0.3	82.9 ± 0.6	66.7 ± 0.7	44.0 ± 0.7	35.3 ± 1.5	62.9
DRO	32.2 ± 3.7	97.9 ± 0.1	77.5 ± 0.1	83.1 ± 0.6	67.1 ± 0.3	42.5 ± 0.2	32.8 ± 0.2	61.8
Mixup	31.2 ± 2.1	98.1 ± 0.1	78.6 ± 0.2	83.7 ± 0.9	68.2 ± 0.3	46.1 ± 1.6	39.4 ± 0.3	63.6
MLDG	36.9 ± 0.2	98.0 ± 0.1	77.1 ± 0.6	82.4 ± 0.7	67.6 ± 0.3	45.8 ± 1.2	42.1 ± 0.1	64.2
CORAL	29.9 ± 2.5	98.1 ± 0.1	77.0 ± 0.5	83.6 ± 0.6	68.6 ± 0.2	48.1 ± 1.3	41.9 ± 0.2	63.9
MMD	42.6 ± 3.0	98.1 ± 0.1	76.7 ± 0.9	82.8 ± 0.3	67.1 ± 0.5	46.3 ± 0.5	39.3 ± 0.9	64.7
DANN	29.0 ± 7.7	89.1 ± 5.5	77.7 ± 0.3	84.0 ± 0.5	65.5 ± 0.1	45.7 ± 0.8	37.5 ± 0.2	61.2
C-DANN	31.1 ± 8.5	96.3 ± 1.0	74.0 ± 1.0	81.7 ± 1.4	64.7 ± 0.4	40.6 ± 1.8	38.7 ± 0.2	61.1

Model selection method: Test-domain validation set (oracle)

Algorithm	CMNIST	RMNIST	VLCS	PACS	Office-Home	TerraInc	DomainNet	Avg
ERM	58.5 ± 0.3	98.1 ± 0.1	77.8 ± 0.3	87.1 ± 0.3	67.1 ± 0.5	52.7 ± 0.2	41.6 ± 0.1	68.9
IRM	70.2 ± 0.2	97.9 ± 0.0	77.1 ± 0.2	84.6 ± 0.5	67.2 ± 0.8	50.9 ± 0.4	36.0 ± 1.6	69.2
DRO	61.2 ± 0.6	98.1 ± 0.0	77.4 ± 0.6	87.2 ± 0.4	67.7 ± 0.4	53.1 ± 0.5	34.0 ± 0.1	68.4
Mixup	58.4 ± 0.2	98.0 ± 0.0	78.7 ± 0.4	86.4 ± 0.2	68.5 ± 0.5	52.9 ± 0.3	40.3 ± 0.3	69.0
MLDG	58.4 ± 0.2	98.0 ± 0.1	77.8 ± 0.4	86.8 ± 0.2	67.4 ± 0.2	52.4 ± 0.3	42.5 ± 0.1	69.1
CORAL	57.6 ± 0.5	98.2 ± 0.0	77.8 ± 0.1	86.9 ± 0.2	68.6 ± 0.4	52.6 ± 0.6	42.1 ± 0.1	69.1
MMD	63.4 ± 0.7	97.9 ± 0.1	78.0 ± 0.4	87.1 ± 0.5	67.0 ± 0.2	52.7 ± 0.2	39.8 ± 0.7	69.4
DANN	58.3 ± 0.2	97.9 ± 0.0	80.1 ± 0.6	85.4 ± 0.7	65.6 ± 0.3	51.6 ± 0.6	38.3 ± 0.1	68.2
C-DANN	62.0 ± 1.1	97.8 ± 0.1	80.2 ± 0.1	85.7 ± 0.3	65.6 ± 0.3	51.0 ± 1.0	38.9 ± 0.1	68.7

Gradient Matching for Domain Generalization


Inter Domain Gradient Matching - IDGM



If there are only x_1 and x_2 types of data, but when $0 \leq f_i \leq 1$ given $\forall f_i \in \mathbb{R}$, will the model distinguish class well only using invariant feature f_1 ?

- **In reality, No.** It depends on **correlation** between class and each feature.
- Many factors can contribute to $Corr(f_i)$, e.g. feature preference, population, etc.
- However, in this toy example, f_i can have only binary value, 0 or 1. So, the model will be use the invariant feature f_1 **“mainly”**.
- If we want model to use **spurious feature** f_2, f_3 , rather than **invariant feature** f_1 , we can inject some **noise** about f_1 .
- Finally, **we can understand that the role of x_3 is the noise for original correlation.**


• Inter Domain Gradient Matching - IDGM



We can expect that the model can learn invariant feature across domain, if we maximize the inner product of gradient for each domain.

$$G_1 = \mathbb{E}_{\mathcal{D}_1} \frac{\partial l((x, y); \theta)}{\partial \theta}, \quad G_2 = \mathbb{E}_{\mathcal{D}_2} \frac{\partial l((x, y); \theta)}{\partial \theta}$$

$$\mathcal{L}_{\text{idgm}} = \mathcal{L}_{\text{erm}}(\mathcal{D}_{tr}; \theta) - \gamma \underbrace{\frac{2}{S(S-1)} \sum_{i,j \in S}^{i \neq j} G_i \cdot G_j}_{\text{GIP, denote as } \hat{G}}$$



- θ at initial point
- ➡ Trajectory of θ during IDGM training
- ➡ Trajectory of θ during ERM training
- - - Gradient of the loss for domain 1
- - - Gradient of the loss for domain 2

$$G_i \cdot G_j > 0$$

$$G_i \cdot G_j < 0$$


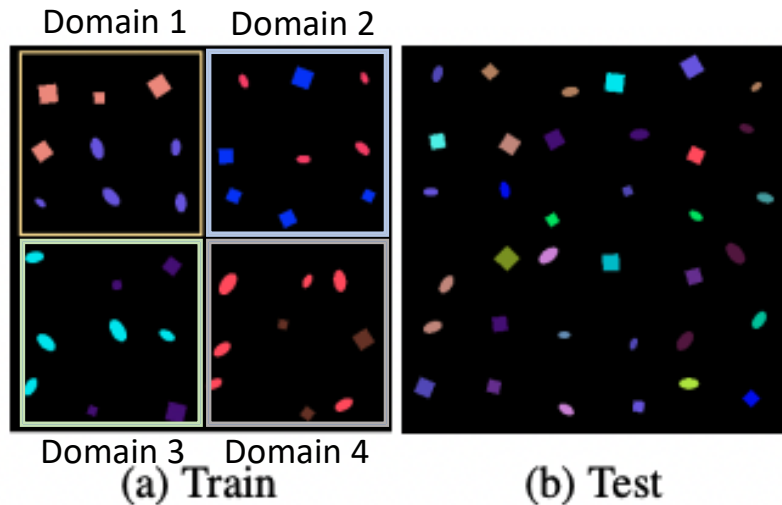


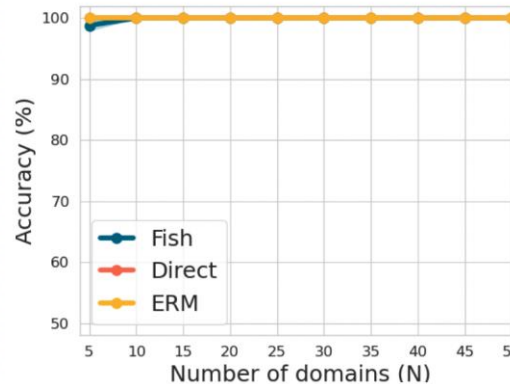

Table 1: Performance comparison on the linear dataset.

Method	train acc.	test acc.	W	b
ERM	97%	57%	[2.8, 3.3, 3.3, 0.0]	-2.7
IDGM	93%	93%	[0.4, 0.2, 0.2, 0.0]	-0.4
Fish	93%	93%	[0.4, 0.2, 0.2, 0.0]	-0.4

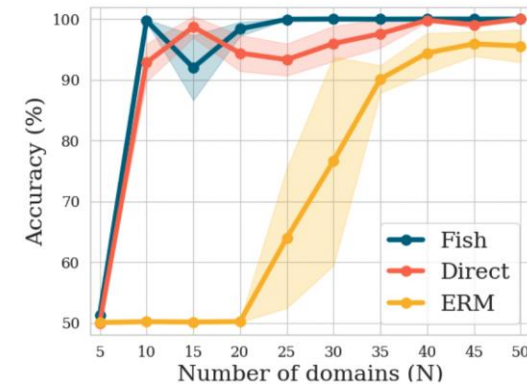
• CDSprites-N Dataset



Label (Invariant) : Shape Attribute (Spurious) : Color



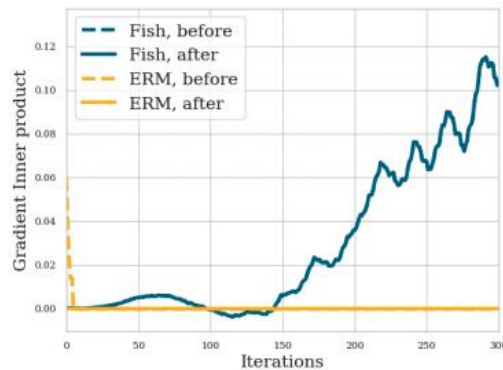
(a) Train




(b) Test

- When the number of domain is small, the model can easily remember the color, so it predict the result by using color instead of shape.
- But as the number of domain increase enough, IDGM, Fish and ERM can learn shape feature, rather than color.
- IDGM and Fish need the small number of domains than ERM. **So, we can say that IDGM and Fish have stronger generalization capabilities.**


Tracking Gradient Inner Product (GIP)



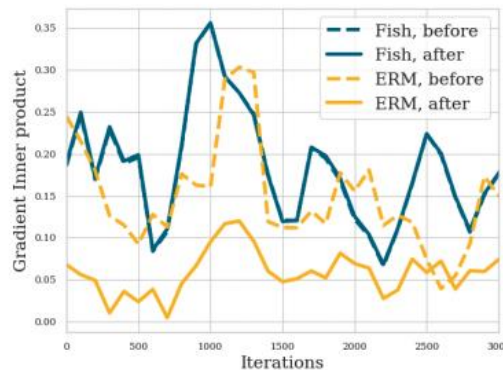
(a) CDSprites-N



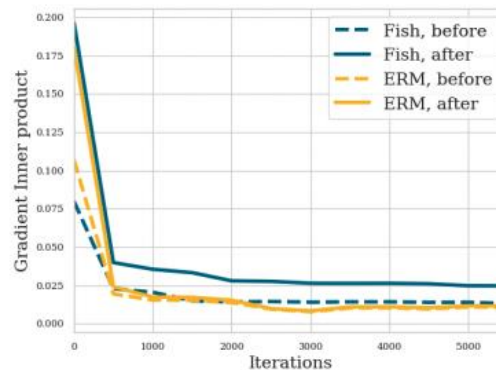
(b) CAMELYON17



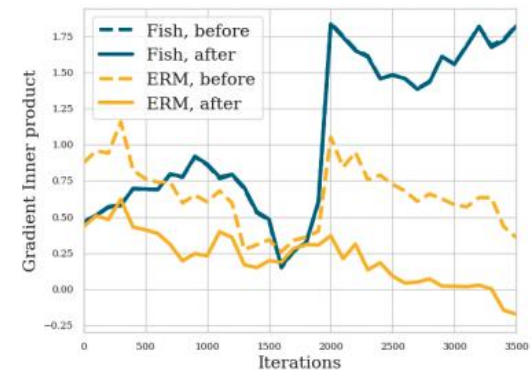
(c) CIVILCOMMENTS



(d) FMoW



(e) POVERTY



(f) IWILDCAM

(dot) Compute GIP before model update / (solid) Compute GIP after model update

Comparing ERM with Fish, ERM often results in the decrease of GIP, while for Fish it can either increase significantly or at least stay at same level.

• IDGM to Fish

Unfortunately, IDGM algorithms is too expensive to compute.

So, we will approximate hessian to first-order optimization. → Fish

Algorithm 1 Fish.

```

1: for iterations = 1, 2, ... do
2:    $\tilde{\theta} \leftarrow \theta$ 
3:   for  $\mathcal{D}_i \in \text{permute}(\{\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_S\})$  do
4:     Sample batch  $d_i \sim \mathcal{D}_i$ 
5:      $\tilde{g}_i = \mathbb{E}_{d_i} \left[ \frac{\partial l((x, y); \tilde{\theta})}{\partial \tilde{\theta}} \right]$  // Grad wrt  $\tilde{\theta}$ 
6:     Update  $\tilde{\theta} \leftarrow \tilde{\theta} - \alpha \tilde{g}_i$ 
7:   end for
8:
9:   Update  $\theta \leftarrow \theta + \epsilon(\tilde{\theta} - \theta)$ 
10: end for
    
```

Conventionally, this type of algorithm is named by vertebrate name.
Fish algorithm is inspired by Reptile algorithm.

Loss of IDGM

$$\mathcal{L}_{\text{idgm}} = \mathcal{L}_{\text{erm}}(\mathcal{D}_{tr}; \theta) - \underbrace{\gamma \frac{2}{S(S-1)} \sum_{i,j \in S}^{i \neq j} G_i \cdot G_j}_{\text{GIP, denote as } \hat{G}}$$

Algorithm 2 Direct optimization of IDGM.

```

1: for iterations = 1, 2, ... do
2:    $\tilde{\theta} \leftarrow \theta$ 
3:   for  $\mathcal{D}_i \in \text{permute}(\{\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_S\})$  do
4:     Sample batch  $d_i \sim \mathcal{D}_i$ 
5:      $g_i = \mathbb{E}_{d_i} \left[ \frac{\partial l((x, y); \theta)}{\partial \theta} \right]$  // Grad wrt  $\theta$ 
6:
7:   end for
8:    $\bar{g} = \frac{1}{S} \sum_{s=1}^S g_s$ ,  $\hat{g} = \frac{2}{S(S-1)} \sum_{i,j \in S}^{i \neq j} g_i \cdot g_j$ 
9:   Update  $\theta \leftarrow \theta - \epsilon(\bar{g} - \gamma(\partial \hat{g} / \partial \theta))$ 
10: end for
    
```

$$\nabla f = \begin{pmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{pmatrix} \quad \text{and} \quad \nabla^2 f = \begin{pmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{pmatrix}$$

Gradient
Hessian

• Key-point in Fish Derivation

Algorithm 1 Fish.

```

1: for iterations = 1, 2, ... do
2:    $\tilde{\theta} \leftarrow \theta$ 
3:   for  $\mathcal{D}_i \in \text{permute}(\{\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_S\})$  do
4:     Sample batch  $d_i \sim \mathcal{D}_i$ 
5:      $\tilde{g}_i = \mathbb{E}_{d_i} \left[ \frac{\partial l((x, y); \tilde{\theta})}{\partial \tilde{\theta}} \right]$  // Grad wrt  $\tilde{\theta}$ 
6:     Update  $\tilde{\theta} \leftarrow \tilde{\theta} - \alpha \tilde{g}_i$ 
7:   end for
8:
9:   Update  $\theta \leftarrow \theta + \epsilon(\tilde{\theta} - \theta)$ 
10: end for
    
```

If meta step $\epsilon = 1$, ERM?

Yes, but it not Fish

Fish Condition

1) Infinite Inner loop

2) Only in small α

If even one condition is violated,
approximation is not established.

Algorithm 2 Direct optimization of IDGM.

```

1: for iterations = 1, 2, ... do
2:    $\tilde{\theta} \leftarrow \theta$ 
3:   for  $\mathcal{D}_i \in \text{permute}(\{\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_S\})$  do
4:     Sample batch  $d_i \sim \mathcal{D}_i$ 
5:      $g_i = \mathbb{E}_{d_i} \left[ \frac{\partial l((x, y); \theta)}{\partial \theta} \right]$  // Grad wrt  $\theta$ 
6:   end for
7:    $\bar{g} = \frac{1}{S} \sum_{s=1}^S g_s$ ,  $\hat{g} = \frac{2}{S(S-1)} \sum_{i,j \in S}^{i \neq j} g_i \cdot g_j$ 
8:   Update  $\theta \leftarrow \theta - \epsilon(\bar{g} - \gamma(\partial \hat{g} / \partial \theta))$ 
10: end for
    
```

$\{d_i\}_{i=1}^S$, where $d_i := \{x_b, y_b\}_{b=1}^B$ denotes a minibatch at step i randomly drawn from one of the available domains in $\{\mathcal{D}_1, \dots, \mathcal{D}_S\}$.

$$\bar{G} = \frac{1}{S} \sum_{s=1}^S G_s$$

$$G_f = \mathbb{E}[(\theta - \tilde{\theta})] - \alpha S \cdot \bar{G}, \text{ Fish update - } \alpha S \cdot \text{ERM grad}$$

$$G_g = -\partial \hat{G} / \partial \theta, \text{ grad of } \max_{\theta}(\hat{G})$$

$$= -\frac{2}{S(S-1)} \sum_{i,j \in S}^{i \neq j} \frac{\partial}{\partial \theta} G_i \cdot G_j$$

$$\tilde{g}_i = \mathbb{E}_{d_i} \left[\frac{\partial l((x, y); \theta_i)}{\partial \theta_i} \right] \text{ (gradient at step } i, \text{ wrt } \theta_i)$$

$$\theta_{i+1} = \theta_i - \alpha \tilde{g}_i \text{ (sequence of parameters)}$$

$$g_i = \mathbb{E}_{d_i} \left[\frac{\partial l((x, y); \theta_1)}{\partial \theta_1} \right] \text{ (gradient at step } i, \text{ wrt } \theta_1)$$

$$H_i = \mathbb{E}_{d_i} \left[\frac{\partial^2 l((x, y); \theta_1)}{\partial \theta_1^2} \right] \text{ (Hessian at initial point)}$$

Perform **Taylor Approximation** for \tilde{g}_i near θ_1 .

$$\tilde{g}_i = l'_i(\theta_i)$$

$$= l'_i(\theta_1) + l''_i(\theta_1)(\theta_i - \theta_1) + \underbrace{\mathcal{O}(\|\theta_i - \theta_1\|^2)}_{=\mathcal{O}(\alpha^2)}$$

$$= g_i + H_i(\theta_i - \theta_1) + \mathcal{O}(\alpha^2)$$

$$= g_i - \alpha H_i \sum_{j=1}^{i-1} \tilde{g}_j + \mathcal{O}(\alpha^2). \text{ Ignore from second-order}$$

$$\tilde{g}_j = g_j + \mathcal{O}(\alpha) \text{ Ignore from first-order}$$

let us consider performing two steps in inner-loop updates, i.e. $S = 2$.

$$\begin{aligned} \theta - \tilde{\theta} &= \alpha(\tilde{g}_1 + \tilde{g}_2) \\ &= \alpha \underbrace{(g_1 + g_2)}_{\textcircled{1}} - \alpha^2 \underbrace{H_2 g_1}_{\textcircled{2}} + \mathcal{O}(\alpha^3) \end{aligned}$$

$$\textcircled{1} = \mathbb{E}_{1,2} [g_1 + g_2] = G_1 + G_2$$

$$\begin{aligned} \textcircled{2} &= \mathbb{E}_{1,2} [H_2 g_1] = \mathbb{E}_{1,2} [H_1 g_2] \text{ (interchanging indices)} \\ &= \frac{1}{2} \mathbb{E}_{1,2} [H_2 g_1 + H_1 g_2] \text{ (averaging last two eqs)} \end{aligned}$$

$$\begin{aligned} &= \frac{1}{2} \mathbb{E}_{1,2} \left[\frac{\partial (g_1 \cdot g_2)}{\partial \theta_1} \right] \\ &= \frac{1}{2} \cdot \frac{\partial (G_1 \cdot G_2)}{\partial \theta_1} \end{aligned}$$

$$\mathbb{E}[\theta - \tilde{\theta}] = \alpha(G_1 + G_2) + \frac{\alpha^2}{2} \cdot \frac{\partial (G_1 \cdot G_2)}{\partial \theta_1} + \mathcal{O}(\alpha^3)$$

We can also expand this to the general case where $S \geq 2$:

$$\begin{aligned} \mathbb{E}[\theta - \tilde{\theta}] &= \alpha \sum_{s=1}^S G_s - \frac{\alpha^2}{S(S-1)} \sum_{i,j \in S}^{i \neq j} \frac{\partial (G_i \cdot G_j)}{\partial \theta_1} + \mathcal{O}(\alpha^3). \\ \bar{G} &= \frac{1}{S} \sum_{s=1}^S G_s. \end{aligned}$$

$$\begin{aligned} G_f &= \mathbb{E}[\theta - \tilde{\theta}] - \alpha S \bar{G} \\ &= -\frac{\alpha^2}{S(S-1)} \sum_{i,j \in S}^{i \neq j} \frac{\partial}{\partial \theta_1} G_i \cdot G_j \end{aligned}$$

$$\lim_{\alpha \rightarrow 0} \frac{G_f \cdot G_g}{\|G_f\| \cdot \|G_g\|} = 1.$$

Cauchy-Schwartz Inequality

$$\therefore G_f \approx k G_g$$

Almost Same Direction

- Domain Generalization
 - Domain Generalization VS Subpopulation Shift
 - Wilds 1.2v Benchmark
 - Algorithms for Domain Generalization
 - DomainBed Benchmark, Model Selection Method
- Gradient Matching for Domain Generalization
 - IDGM – Inter Domain Gradient Matching
 - IDGM and Fish Algorithms
 - Derivation IDGM → Fish